

jbi training

C++ 20 and 23

Programming C++ 20 and C++ 23

Contents	Slide
Programming C++ 20	1
About This Course	2
Template Features	9
Concepts	19
Constraints and Concepts	37
Span	57
Ranges and Views	62
Threading	97
Atomic	114
General C++ 20 Features	120
General C++ 20 Features Continued	139
Spaceship Operator	140
Format Library	163
Print Library (C++ 23)	177
C++ 23 Operators	182
C++ 23 Fold Algorithms	188
C++ Monad	194
Appendix 1	207
Modules	208
Appendix 2	226
Coroutines	227

Programming C++20

About This Course

- This course overviews some recent C++ 20 features
- In the development of the C++ language there have been some significant standard releases and some less significant
- The major standard releases (since 1998) are:
 - C++ 11 (R value reference and move)
 - C++ 20

IDE Support

- This course overviews some recent C++ features
 - Visual Studio and GNU History:
 - Visual Studio 2008 – C++98
 - Visual Studio 2010 – C++0x (partial)
 - Visual Studio 2012/2013 – C++11 (partial)
 - Visual Studio 2015 – C++14 (partial)
 - Visual Studio 2017 – C++17 (partial)
 - Visual Studio 2019 – C++20 (partial)
 - Visual Studio 2022 – C++23 (partial?)
 - GNU C++ (depends on version!)
 - The support also depends upon the update

All Trademarks acknowledged: Microsoft Visual Studio .NET, Visual C++, C#, VB.NET, SQL Server, Windows

©Copyright Allotropical Ltd. 1998-2024 All rights reserved.

3

Compiler Support

- The detailed support for C++ is a moving target, see the page on ‘cppreference’

https://en.cppreference.com/w/cpp/compiler_support

- Within Visual Studio use the project properties language options to select C++20
 - For some features need to choose latest!
- For GCC compilers using options:

`-std=c++20 -fcoroutines`

Additional flag for coroutines

©Copyright Allotropical Ltd. 1998-2024 All rights reserved.

4

Feature Macros

- Many defines which can be checked to determine if the feature is supported:

```
#include <version>
int main()
{
    #ifdef __cpp_lib_barrier
    std::cout << "Barrier supported" << std::endl;
    #endif
    #ifdef __cpp_lib_bind_front
    std::cout << "Bind front supported" << std::endl;
    #endif
    #ifdef __cpp_using_enum
    std::cout << "Using enum supported" << std::endl;
    #endif
    ...
}
```

If project property set
for language C++20

Barrier supported
Bind front supported
Using enum supported
...

C++17 Features

- The C++17 standard include a number of features considered useful:
 - Nested Namespaces
 - noexcept part of type system
 - Template Argument Deduction
 - Initialisation within if and switch!
 - Constexpr If
 - Parallel algorithms

C++17 Functions and Types

- The C++17 include a number of useful functions and types:
 - Structured Bindings
 - std::invoke, std::apply and std::make_from_tuple
 - Inline Static Variables
 - Copy Elision and RVO
 - Fold Expressions
 - std::string_view
 - std::optional, std::variant
 - std::byte
 - File System

C++20 Benefits

- The benefits of moving to C++ 20 are many and fall into a number categories:
 - Efficiency
 - Constexpr, consteval and constinit
 - Likely and Unlikely attributes
 - Compile time diagnostics
 - Concepts
 - Code simplification
 - Abbreviated template syntax
 - Lambdas
 - Operator definition (spaceship)
 - Conditional explicit

Template Features

- This section is an overview of the new template related features:
 - Abbreviated Template Syntax
 - Lambda Template Syntax
 - Template Parameters
 - Class Non-type Template Parameters
 - Non-type Parameter Pointer to Function
 - Variadic Templates

Abbreviated Templates Syntax

- C++20 supports a simplified template syntax
- Old style template function:

```
template<typename T>
T old_doubler(T t) { return t + t; }
```
- Simplified variation:

```
auto doubler(auto n) { return n + n; }
```
- Can also use concepts:

```
auto doubler_numeric(Numeric auto n) { return n + n; }
```

Multiple Template Arguments

- C++20 supports a simplified template syntax
- Old style template function:

```
template<typename T>
T my_max(T x, T y) { return x < y ? y : x; }
```

- What if the calling types are different? How do we define the return type?

```
template<typename T, typename U>
auto my_max(T x, U y) -> decltype(x+y)
    { return x < y ? y : x; }
```

Expression to allow deduction of type?

- Simplified:

```
auto amy_max(auto x, auto y) { return x < y ? y : x; }
```

Lambda Template Syntax

- A generic lambda does not easily convey type information (use of auto):

```
auto nested_type = [](auto t)
{
    return typeid(std::declval<decltype(t)::value_type>).name();
};

std::cout << nested_type(std::vector{1, 2}) << std::endl;
```

Requires use of decltype
to deduce type

```
auto nested_type = []<typename T>(T t)
{
    return typeid(T::value_type).name();
};

std::cout << nested_type(std::vector{1, 2}) << std::endl;
```

Type parameter, can
also use concepts

Template Parameters

- C++17 introduced the ability to have non-type template parameter deduction
 - Non-type parameters essentially restricted to integral types
- C++ 20 extends non-type parameters to floating point types and class types
 - Class types should be structural types (simple!)
 - Class and base classes have public members
 - Members are structural types!

C++17 Template Non-type Deduction

```
int main()
{
    SomeData<int,7u> sd;
    auto a = sd.getVal(4);
    std::cout << a << std::endl;
}

#include <iostream>
#include <numeric>
template <typename T, auto Ndecltype(N) i) const { return data[i]; }
};
```

Deduce non-type parameter

Deduce parameter

Class Non-type Template Parameter

```
Requires constexpr  
constructor  
  
Class Type  
  
template<DataClass DC>  
int dc_calc( int b,int c ) {  
    return DC.calc(b,c);  
}  
int main(){  
  
    std::cout << "Result: " << dc_calc<DataClass{10}>(1,3) << std::endl;  
}
```

struct DataClass {
 int a;
 constexpr DataClass(int n) :a{ n } {}

 int calc(int b, int c)const { return a - b + c; }
};

Literal Object

Lambda as Non-type Parameter

- Non-type parameter supports pointer to function
 - This can now accept a lambda expression

```
using FunPtr = double(double);  
  
template<FunPtr F>  
class Command {  
public:  
    double execute(double d) { return F(d); }  
};
```

```
Command<[](double d)->double {return d * 7; }>
```

Constexpr Lambdas

- Non-type parameter supports pointer to function
 - This can now accept a lambda expression

```
using FunPtr = double(double);  
  
template<FunPtr F>  
class Command {  
public:  
    constexpr double execute(double d) { return F(d); }  
};
```

```
Command<[](double d) constexpr ->double {return d * 7; }>
```

Variadic Templates

- C++20 supports capture of variadic arguments within a lambda capture:

```
template<typename... Args> Variadic type pack  
auto lambda_fun_outer = [](Args&&... args) { Variadic argument pack  
    Capture and forwarding  
    of argument pack  
    return ([... args = std::forward<Args>(args)]() {return (args + ...); })();  
};
```

Concepts

- This section gives and overview of motivation and using concepts:
 - Concepts Introduction
 - Named Requirements
 - Named Requirements and Type Traits
 - Template Problems
 - Templates and SFINAE
 - Templates and concepts

Concepts Introduction

- This section gives an introduction to the motivation for the use of concepts
- The idea of Named Requirements will be introduced before considering some of the issues with working with templates
- Existing solutions will be considered using pre C++20 features

Named Requirements

- Named requirements have been used within the C++ Standard for many years
 - They define requirements on types for template instantiation
 - Not meeting these requirements can result in complex compiler messages
 - Some of these requirements can be checked for using type traits
- It has been largely down to the developer to meet these requirements

Named Requirements

- The expected behaviour of the Standard Library is largely defined in terms of ‘Named Requirements’
- There are now many named requirements, which fall into a number of categories:
 - Basic requirements
 - Container
 - Iterator
 - Concurrency

Basic Named Requirements

- These ‘Named Requirements’ provide for common features:

Requirement	Description (fairly self explanatory!)
DefaultConstructible	Object can be default constructed
CopyConstructible	lvalue constructable
MoveConstructible	rvalue constructable
CopyAssignable	Assignable from lvalue
MoveAssignable	Assignable from rvalue
Destructible	Object can be destroyed

Named Requirements Checking

- Named requirements are capable of being checked through the use of type traits, e.g. DefaultConstructible:
 - The three type traits below express a number of variations on default constructable:

Type trait (C++11)	Addition in C++17
is_default_constructible	is_default_constructible_v
is_trivially_default_constructible	is_trivially_default_constructible_v
is_nothrow_default_constructible	is_nothrow_default_constructible_v

type traits

- Type traits were introduced with C++11
 - There are a wide range type traits
 - Additional type traits add with C++17 to simplify usage, e.g.

```
template<typename T>
    inline constexpr bool is_default_constructible_v =
        is_default_constructible<T>::value;
```

- These are now being used to formalize ‘Named Requirements’ within the C++20 standard

type_traits

- Type_traits allow determining characteristics for types/functions etc.

Type traits	value
is_abstract	True if abstract class
is_array	True if array
is_pod	True if plain old data (trivial and standard layout)
is_assignable	True if assignment possible between types
is_base_of	True if base of another type
is_class	True if class (or struct)
is_function	True if type is function

type_traits (trivial)

- Type_traits allow determining characteristics for types/functions etc.

Type traits
is_trivial
is_trivially_copyable
is_trivially_assignable
is_trivially_copy_assignable

- Each characteristic has a number of requirements
 - Largely regarding only implicitly supplied constructors; assignment operators and destructor
 - Also types not containing virtual member functions

Named Requirements - Categories

- There are now many named requirements:

Category	Examples
Basic	DefaultConstructable, CopyConstructable, CopyAssignable
Type	TriviallyCopyable, PODType
Container	Container, SequentialContainer, AssociativeContainer
ContainerElement	DefaultInsertable, CopyInsertable
Iterator	LegacyIterator, LegacyBidirectionalIterator
Stream I/O	FormattedInputFunction, FormattedOutputFunction
Concurrency	Lockable, SharedLockable, Mutex
Library	EqualityComparable, Allocator, Swappable
Other	NumericType, LiteralType, BitmaskType
Random Number	SeedSequence, RandomNumberEngine

Name Requirements and Type Traits

- Some of the Named Requirements are associated with type traits
 - There are many predefined type traits within the header file ‘type_traits’
- These type traits can be used in conjunction with ‘static_assert’ to give a compiler message:

```
static_assert(std::is_constructible<A>::value,  
             "A not constructible!");
```

C++ 11
- ```
static_assert(std::is_constructible_v<A>,
 "A not constructible!");
```

C++ 17

# Multiple Type Traits

- Multiple type traits can be applied using conjunction or disjunction
  - Logical ‘and’:

```
std::conjunction_v<...>
```
  - Logical ‘or’:

```
std::disjunction_v<...>
```
- These type traits can be useful, especially within variadic templates

# Template Problems

- A simple template function could be instantiated with a wide range of types
  - Some types may not be appropriate!

Template instantiated for any type with a '<' operator

Works for 'int', but would not work as expected for const char \*

```
template<typename T> inline T my_max(T x, T y)
{
 return x < y ? y : x;
}
int main()
{
 int a = 65, b = 82, result;
 result = my_max(a, b);
 ...
}
```

# Use of static\_assert

- Static assert can be used cause compiler error of unexpected type:

```
template<typename T> inline T my_max(T x, T y)
{
 static_assert(std::is_integral_v<T> || std::is_floating_point_v<T>, "Not a number!");
 return x < y ? y : x;
}

int main()
{
 const char* pch = "Hello", *pch2 = "Greetings", *result;
 result = my_max(pch, pch2); // Compiler error due to attempted use of const char *
}
```

# Templates and SFINAE

- A different alternative is to use SFINAE to prevent instantiation!

```
template<typename T> inline std::enable_if_t<
 std::is_integral_v<T> ||
 std::is_floating_point_v<T>, T>
 my_max(T x, T y)
{
 return x < y ? y : x;
}
```

Multiple type traits

- Alternative versions can be provided again using ‘enable\_if’

## Templates and SFINAE continued...

- Alternative versions of the template function can be provided where the types satisfy different traits!

```
template<typename T> inline std::enable_if_t<
 std::is_convertible_v<T, const char*>, T>
 my_max(T x, T y)
{
 return std::string(x) < std::string(y) ? y : x;
}
```

## Templates and SFINAE continued...

- The previous slide showed the use of ‘enable\_if’ in the return type, but it could appear in the template parameters:

```
template<typename T, typename =
 std::enable_if_t<std::is_convertible_v<T, const char*>>>
inline T my_max(T x, T y)
{
 return std::string(x) < std::string(y) ? y : x;
}
```

- However, this approach is not ideal!

## Templates and Concepts

- The C++ 20 alternative is to use concepts
  - A max template function just for floating point types:

Concept to required type

```
template<std::floating_point T>
inline T my_max(T x, T y)
{
 return x < y ? y : x;
}
```

# Constraints and concepts

- This section gives an overview of Constraints and concepts:
  - Constraints using requires
  - Concepts
  - Predefined Concepts
  - Defining concepts
  - Use of Concepts
  - Concepts and requires
  - Defining Concepts

## Constraints using requires

- Constraints can be applied to template parameters using requires expressions:

```
template<typename T>
 requires(std::is_integral_v<T> || std::is_floating_point_v<T>)
 inline T a_max(T x, T y) {
 return x < y ? y : x;
 }
```

Requires expression

Disjunction or conjunction (&&) can be used between type traits

# Concepts

- C++ templates provide a very flexible means of implementing generic programming
  - Instantiation of templates, with inappropriate type, can result in all manner of errors!
- The use of ‘concepts’ can help with the definition of templates in order to restrict instantiation to viable options

## Predefined Concepts

- There are many predefined concepts within the ‘concept’ header file:

| Concept categories | Examples                                                |
|--------------------|---------------------------------------------------------|
| Core language      | same_as, integral, floating_point, destructible         |
| Comparison         | equality_comparable, totally_ordered                    |
| Object             | movable, copyable, regular                              |
| Callable           | invocable, regular_invocable, predicate,                |
| Iterator           | input_iterator, output_iterator, random_access_iterator |
| Algorithm          | indirect_unary_predicate, indirect_binary_predicate     |
| Ranges             | range, view, input_range, random_access_range           |

# Useful Concepts

- These are a few concepts in more details:

| Concept                       | Description                                                     |
|-------------------------------|-----------------------------------------------------------------|
| std::same_as<T>               | Satisfied if the type is the same as T                          |
| std::convertible_to<T>        | Satisfied if the type can be converted to T                     |
| std::invocable<F, ...>        | Satisfied if the type is callable with the arguments            |
| std::regular_invocable<F,...> | Std::invocable, plus, does not change function or arguments     |
| std::predicate<F, ...>        | Satisfied if the type is regular_invocable and boolean-testable |
| std::integral                 | Satisfied if the type is integral                               |
| std::floating_point           | Satisfied if the type is floating point                         |

## std::invocable<>

- Require to put an explicit constraint on callable passed to template:

```
template <typename Func, typename T>
requires std::invocable<Func, T, T> // constraint
T operate(Func func, T a, T b) {
 return func(a,b);
}

int main() {
 auto mult = [](int x, int y) {return x*y;};
 auto result = operate(mult, 5, 7);
 std::cout << result << std::endl;
 return 0;
}
```

# Defining concepts

- Concepts can be defined using template syntax:

```
template<typename T>
concept Integral = std::is_integral<T>::value;
template<typename T>
concept Floating = std::is_floating_point<T>::value;
```

Concept defined in terms of type traits

Use of concept to provide constraint

```
template<Integral T>void do_output(T t) {
 std::cout << "Integral: " << t << std::endl;
}
template<Floating T>void do_output(T t) {
 std::cout << "Floating: " << t << std::endl;
}
```

## Use of Concepts

- The previous slide used concepts in place of typename within template definitions
- Another way of using a concept is as a requirement using requires:

```
template<typename T>void do_output2(T t)requires Integral<T> {
 std::cout << t << std::endl;
}
template<typename T>void do_output2(T t)requires Floating<T> {
 std::cout << t << std::endl;
}
```

# Composite Concepts

- A composite concept can be defined using multiple type traits or other concepts:

```
template<typename T>
concept Numeric = std::is_integral_v<T>
Disjunction || std::is_floating_point_v<T>;
```

```
template<Numeric T>
inline T square(T n) { return n * n; }
int main()
{
 int a = 2;

 auto result = square(a);
 ...
}
```

Concept restricts type to be used with ‘square’ function

Alternatively:

```
template<typename T>
concept Numeric = Integral<T>
|| Floating<T>;
```

# Concepts and Requires

- An alternative way of using constraints is by using requires:

```
template<typename T>
inline T square(T n) requires Numeric<T> { return n * n; }
```

```
int main()
{
 int a = 2;

 auto result = square(a);
 ...
}
```

Applying concept through requires

# Defining concepts

- Concepts can separate out requirements, e.g. incrementing:

```
template<typename It>
concept MyIncrementor = requires(It it) {
 {++it} -> std::same_as<It&>;
 {it++} -> std::same_as<It>;
};
```

Support both prefix and  
postfix incrementation,  
with return specification

- This concept can be used in its own right or combined with others to form more complex concepts

# Defining concepts continued...

- Concepts can separate out requirements, e.g. dereferencing:

```
template<typename It>
concept MyDerefer = requires(It it) {
 {*it}->std::convertible_to<typename It::reference>;
};
```

Requires dereference  
operator \*

Concept for return type

# Compose concepts

- The two previous concepts of MyIncrementor and MyDerefer:

```
template<typename It>
concept MyIterator = MyIncrementor<It> && MyDerefer<It>;
```

```
template<MyIterator It>
void use_iterators(It begin, It end) {
 for (auto it = begin; it != end; ++it) {
 std::cout << *it << ", ";
 }
 std::cout << std::endl;
}
```

# Full Iterator!

- A more complete implementation of an iterator concept would be of the form:

```
template<typename It >
concept MyIterator2 = requires(It it1, It it2) {
 {++it1} -> std::same_as<It&>;
 {it1++} -> std::same_as<It>;
 {*it1} -> std::convertible_to<typename It::reference>;
 {it1 == it2} -> std::convertible_to<bool>;
 {it1 != it2} -> std::convertible_to<bool>;
};
```

Annotations for the MyIterator2 requirements:

- Multiple variables: Points to the requirement `{++it1} -> std::same_as<It&>;`
- Increment: Points to the requirement `{it1++} -> std::same_as<It>;`
- Dereference: Points to the requirement `{*it1} -> std::convertible_to<typename It::reference>;`
- Comparison: Points to the requirements `{it1 == it2} -> std::convertible_to<bool>;` and `{it1 != it2} -> std::convertible_to<bool>;`

# Interface like concept

- Concepts can be used a little like interfaces (which are available in many languages):

```
template<typename T>
concept Stringable = requires(T t) {
 { t.to_string() } -> std::same_as<std::string>;
};
```

Requirement for  
member function

Requirement for  
return type

## noexcept and const

- The previous slide considered the ability to require a member function signature
- It can be useful to add additional requirements for noexcept or const:

```
template<typename T>
concept Stringable = requires(T t) {
 {std::as_const(t).to_string() } -> std::same_as<std::string>;
} && noexcept(std::declval<const T>().to_string());
```

noexcept operator

Requirement a const object

# Implementing class

- A class satisfying a ‘constraint’ is defined completely independently of the ‘concepts’:

```
class PassengerDetails {
 std::string _name;
 int _weight;
public:
 PassengerDetails(const std::string& name, const int& weight):
 _name(name), _weight(weight){}

 std::string to_string() const { return "Name: " + _name +
 ", baggage weight: " + std::to_string(_weight); }
}
```

# Using Implementing Class

- Implement template function using concept:

```
template<Stringable T>
void do_output(T t) {
 std::cout << t.to_string() << std::endl;
}
```

Tooling within some development environments show member functions available

- Calling function using an object of a class meeting requirement:

```
{
 PassengerDetails pd{ "Fred", 43 };

 do_output(pd);
}
```

# Concept Operator Requirements

- Operators can be represented using concepts:

```
template<typename T>
concept Addable = requires(T n1, T n2) {
 { n1 += n2 } -> std::convertible_to<T&>;
 { n1 + n2 } -> std::convertible_to<T>;
};

template<Addable A>
void calc(A a1, A a2) {
 a1 += 7;

 auto result = a1 + a2;
}
```

Multiple variables

Requirement for operators

Use of concept

# Implementing Operators

- The Adder concept defined satisfied by many existing types, i.e. int, std::string or a custom type:

```
class Number {
 int _val;
public:
 Number(int val):_val(val){}
 Number& operator+=(const Number& n) { _val += n._val; return *this; }
};

inline Number operator+(const Number& lhs, const Number& rhs) {
 Number temp(lhs);
 return temp += rhs;
}
```

# Span

- This section gives an overview of spans:
  - Span Introduction
  - Span Member Functions
  - Span Example
  - Subspan Example

## Span Introduction

- A span represents a view onto a contiguous sequence (does not own the data)
  - Can be initialised with arrays, many sequential containers and ranges
  - Changes may be made through the span
  - Subspans can be created
  - Useful for passing arrays to functions

# Span Member Functions

- There are many member functions which allow a span to be used like a container:

| Member function | Description                                                   |
|-----------------|---------------------------------------------------------------|
| begin, rbegin   | Returns beginning or reverse beginning iterator, respectively |
| end, rend       | Returns end or reverse end iterator, respectively             |
| front, back     | Returns values from front or back, respectively               |
| operator[]      | Index into the span (may throw out_of_range exception)        |
| data            | Returns pointer to start of contiguous data                   |
| size            | Size of the span                                              |
| first, last     | Return spans of given size, starting at the beginning or end  |
| subspan         | Returns a span with given start point and count               |

# Span Example

- Using span as a parameter:

```
inline int mysum(std::span<int> data) {
 int result{};
 for (size_t i = 0; i < data.size(); i++)
 {
 result += data[i];
 }
 return result;
}
int main(){
 int anarray[]{1,2,3,4,5,6,7,8,9};
 int result = mysum(anarray);
 std::cout << result << std::endl;
}
```

The diagram illustrates the use of the `mysum` function with a `std::span` parameter. Annotations explain the following concepts:

- A callout points to the parameter `data` in the `mysum` function definition with the text "Take a span as a parameter".
- A callout points to the `size()` method call in the loop condition with the text "Size can be obtained from span".
- A callout points to the index expression `data[i]` with the text "Index into span".
- A callout points to the declaration of the `anarray` array with the text "'C' style array".
- A callout points to the argument passed to `mysum` with the text "Pass array to be viewed through a span".

# Subspan Example

- Subspans can be created from a span:

```
int anarray[] { 1,2,3,4,5,6,7,8,9 };
std::span<int> sp{ anarray };
Subspan created
from start and count
auto first_part{ sp.subspan(0, 4) };
auto second_part{ sp.subspan(4, 5) };
sum of the separate subspans
int result_from_parts = mysum(first_part) + mysum(second_part);
Sum from span
int result = mysum(sp);
Same results
std::cout << result << " = " << result_from_parts << std::endl;
```

# Ranges and Views

- This section gives an overview of Ranges and Views:
  - Ranges and Views Introduction
  - Ranges
  - Algorithms – Ranges
  - Range Factories
  - Pipelines
  - Range Adapters

# Ranges and Views Introduction

- A new set of features which have been added to the standard library are ranges and views
  - Make use of a number of concepts
- The need for ranges is partly motivated by a problem with working with iterators
  - Algorithms typically require both begin and end iterator to be specified (of the same type)
  - No compiler check that begin and end refer to the same container! Thus, possible runtime error!

## Ranges

- Ranges are generalization to the iterator and algorithm library
- Ranges are an abstraction for:
  - Iterators - [begin, end) – Half open interval
  - Counted sequence – [start, size)
  - Conditional – [start, predicate)
  - Unbounded – [start, ...)

# Range Abstractions

- Examples of abstractions:

| Abstraction | Adapter    |
|-------------|------------|
| Counted     | counted    |
| Conditional | take_while |
| Unbounded   | iota_view  |

# Ranges

- The std::ranges namespace has many functions, including algorithms:

| Function                          | Description                                        |
|-----------------------------------|----------------------------------------------------|
| begin, cbegin,<br>rbegin, crbegin | Returns the various begin iterator                 |
| end, cend, rend,<br>crend         | Returns the various end iterator (sentinel!)       |
| size, ssize                       | Returns unsigned and signed size                   |
| empty                             | Returns Boolean indicating if range is empty       |
| data, cdata                       | Returns pointer to beginning of ‘contiguous’ range |

# Constrained Algorithms

- Constrained algorithms are overloads of algorithms which take a ‘range’ instead of the traditional pair of iterators, i.e. single range or iterator sentinel pair
  - The ‘range’ concept provides for the begin and end iterator
  - The object implementing ‘range’ can simply be passed as one parameter

## Algorithms - Ranges

- Including `<algorithm>` header file now provides algorithms which support the use of ranges:

```
#include <algorithm>
int main()
{
 std::vector<int> data(10);
 std::fill(data.begin(), data.end(), 6);
 std::ranges::fill(data, 4);
}
```

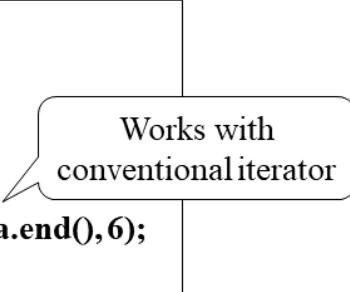
The diagram shows two code snippets side-by-side. The first snippet uses the standard `std::fill` function from the `<algorithm>` header, passing it the `begin()` and `end()` iterators of a `std::vector` and the value 6. A callout bubble points to this line with the text "Traditional algorithm using iterator". The second snippet uses the newer `std::ranges::fill` function from the `<ranges>` header, passing it a range (indicated by curly braces) and the value 4. A callout bubble points to this line with the text "New Algorithm using Ranges".

- Equivalent numeric algorithms do not currently support Ranges

# Algorithms - Ranges

- The previous slide showed the use of the `std::ranges::fill` with a range, however there is an overload which uses an iterator and ‘sentinel’
  - The idea of a ‘Sentinel’ is an end only iterator!

```
#include<algorithm>
int main()
{
 std::vector<int> data(10);
 std::ranges::fill(data.begin(), data.end(), 6);
}
```

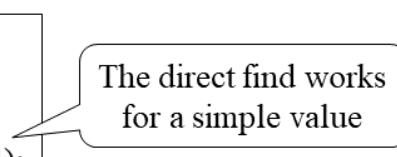


Works with conventional iterator

# Algorithms - Projections

- Another feature of the new ranges algorithms is the optional use of projections
- Consider the find algorithm:

```
{
 std::vector data{ 2,4,6,8 };
 auto found = std::ranges::find(data, 4);
 std::cout << *found << std::endl;
}
```



The direct find works for a simple value

- What if the container stored objects and comparison was required with a member?

# Algorithms - Projections

- Consider the struct:

```
struct Info {
 int val;
 Info(int val):val(val){}
};

{
 std::vector<Info> data{ 2,4,6,8 };

 auto result = std::ranges::find(data, 6, &T::val);

 std::cout << result->val << std::endl;
}
```

Projection for data member

- The ability to ‘invoke’ using the address of a data member is possible using the ‘invoke’ function (C++ 17)

# Sentinel

- The problem which motivated the introduction of the concept of a sentinel was that traditional iteration requires iterators of the same type and of a known range!
  - Iterators work well with containers
- Sentinels allow a generalization of the iterator
  - Allows specialised collections or ranges can be define

# Custom Sentinel Example

- Support a sequence was required which had an unusual ‘end’ requirement
  - Array of values, where -1 denotes the ‘end’
- A traditional iterator could not detect this
  - Therefore define a sentinel

```
struct ArraySentinel {};
```

Illustrates minimal implementation

```
template<typename T>
constexpr bool operator==(T ptr, ArraySentinel) noexcept {
 return *ptr == -1; // Sentinel condition
}
```

# Example Sentinel Usage

- Sentinel used like an ‘end’ iterator:

```
std::array<int, 6> numbers = { 1, 2, 3, 4, 5, -1 };

auto sentinel = ArraySentinel{}; // Create an instance

for (auto it = std::ranges::begin(numbers);
 it != sentinel; ++it) {
 std::cout << *it << " ";
}
std::cout << std::endl;
std::ranges::copy(std::ranges::cbegin(numbers), sentinel,
 std::ostream_iterator<int>(std::cout, ", "));
std::cout << std::endl;
```

! = works!

Used in ranges algorithm

# Ranges Concepts

- Many concepts are defined for working with ranges (below are a number of these concepts):

| Concept                            | Description                                                     |
|------------------------------------|-----------------------------------------------------------------|
| range                              | Requires a begin and end member function                        |
| sized_range                        | Known sized range                                               |
| view                               | View is a range with copy, move and assignment                  |
| input_range,<br>output_range, etc. | Various ranges corresponding to the various iterator categories |
| viewable_range                     | Range can be converted to a view                                |

## An Aside – ‘iota’

- The iota function is used to create an increasing sequence
- The numeric algorithm may be used as:

```
#include <numeric>
```

```
std::vector<int> data(10);
std::iota(data.begin(), data.end(), 1);
```

Populates the vector with  
the values 1 through to 10

- Named after the APL function ‘i’ (index generator)
- The term ‘iota’ now appears in various other places involved with generation of sequences

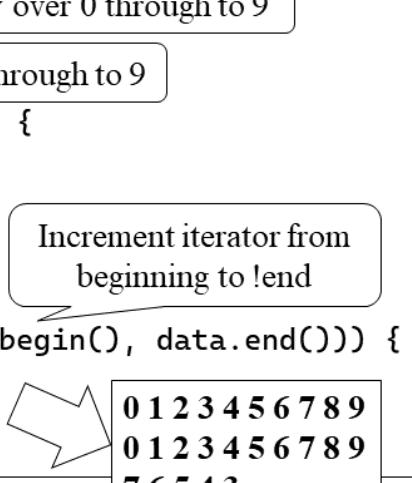
# iota views

- Variations on use of ‘iota’

```
for (int i : std::ranges::iota_view{ 0, 10 }) {
 std::cout << i << ' '; } Create a view over 0 through to 9
std::cout << std::endl; view over 0 through to 9

for (int i : std::views::iota(0, 10)) {
 std::cout << i << ' '; }
std::cout << std::endl; Increment iterator from
beginning to !end

std::vector<int> data{7, 6, 5, 4, 3};
for (auto it : std::views::iota(data.begin(), data.end())) {
 std::cout << *it << ' '; }
std::cout << std::endl;
```



©Copyright Allotropical Ltd. 1998-2024 All rights reserved.

77

# Range Factories

- Various range factories are available:

| Range factory                                | Description                       |
|----------------------------------------------|-----------------------------------|
| ranges::empty_view<br>views::empty           | View contains no elements         |
| ranges::single_view<br>views::single         | View with a single element        |
| ranges::iota_view<br>views::iota             | View with increasing elements     |
| ranges::basic_istream_view<br>views::istream | Elements input from ‘>>’ operator |

©Copyright Allotropical Ltd. 1998-2024 All rights reserved.

78

## istream\_view from keyboard

- Input from many source can be obtained using ‘istream\_view’:

```
auto results = std::ranges::istream_view<std::string>(std::cin);

std::ostream_iterator<std::string> osi(std::cout, " ");
std::ranges::copy(results, osi);
```

- Keyboard input will continue until ‘Ctrl Z’ on Windows and ‘Ctrl D’ on Linux

## istream\_view from std::string

- Values can be read from a string:

```
std::istringstream msg{ "12.3 6.3 7.3" };
auto results = std::ranges::istream_view<double>(msg);

std::ostream_iterator<double> osi(std::cout, " - ");
std::ranges::copy(results, osi);
```

- Output: **12.3 - 6.3 - 7.3 -**

# Dangling

- Care needs to be taken when using ranges, to avoid ‘dangling’:

```
int main()
{
 Found an iterator? get_data() returns a temporary
 auto found = std::ranges::find(get_data(), 3);

 std::cout << *found << std::endl;
}
```

\*found causes compilation error due  
to found not having an operator\*

# Dangling Solution

- To solve the previous problem create an object with appropriate lifetime:

```
auto get_data() { return std::vector{ 1,2,3,4,5 }; }

int main()
{
 auto data = get_data(); Object with larger scope
 auto found = std::ranges::find(data, 3);

 std::cout << *found << std::endl;
}

Dereferencing now safe as
'data' in scope
```

# Dangling

- Some types resolve the problem of dangling, by having borrowed ranges, such as `std::string_view` and `std::span`:

```
using std::string_literals::operator"" s;
{
 auto str{ "Greetings"s };
 auto found = std::ranges::find(std::string_view{str}, 't');

 std::cout << *found << std::endl;
}
```

This is valid as the underlying  
'string' still exists!

Temporary string\_view!

# Pipelines

- Ranges and Views can be used to create pipelines
- The use of multiple traditional algorithms will result in iterating over possibly the same container(s) multiple times
- The use of Pipelines allows a single iteration whilst processing the data
  - R values cannot be used in a pipeline
  - Views can be used within a pipeline

# | operator

- To achieve pipelining the overloaded or ‘|’ operator is used
- The value(s) on the left of the operator are fed into the expression on the right of the operator

```
int data[] = {1,2,3,4,5};
auto results = data
 | std::views::filter([](auto n)
 {return n % 2 == 0;});
```

Range Adapter

Pipeline operator

# Range Adaptors

- Range Adaptors provide various functionality
- Range Adaptors only work with viewable ranges, these include:

| Adaptor    | Returned view   | Description                                          |
|------------|-----------------|------------------------------------------------------|
| filter     | filter_view     | Filters elements based on a predicate                |
| transform  | transform_view  | Transforms elements based on a ‘callable’            |
| take       | take_view       | First set of elements                                |
| take_while | take_while_view | First set of elements whilst predicate true          |
| drop       | drop_view       | Skip the first set of elements                       |
| drop_while | drop_while_view | Skip the first set of elements whilst predicate true |

# Range Adaptors

- Range Adaptors provide various functionality
- Range Adaptors only work with viewable ranges, these include:

| Adaptor | Returned view | Description                |
|---------|---------------|----------------------------|
| join    | join_view     | Flatten the sequence       |
| reverse | reverse_view  | Reverse elements           |
| split   | split_view    | Split sequence of elements |

# Pipelining (lazy evaluation)

- Sequences of data can be processed by pipelining, using range adaptors:

```
auto data = std::vector<int>(20);
std::iota(data.begin(), data.end(), 1);

auto results = data
 | std::views::filter([](auto n) { return n % 2 != 0; });

std::cout << "Odd numbers: ";

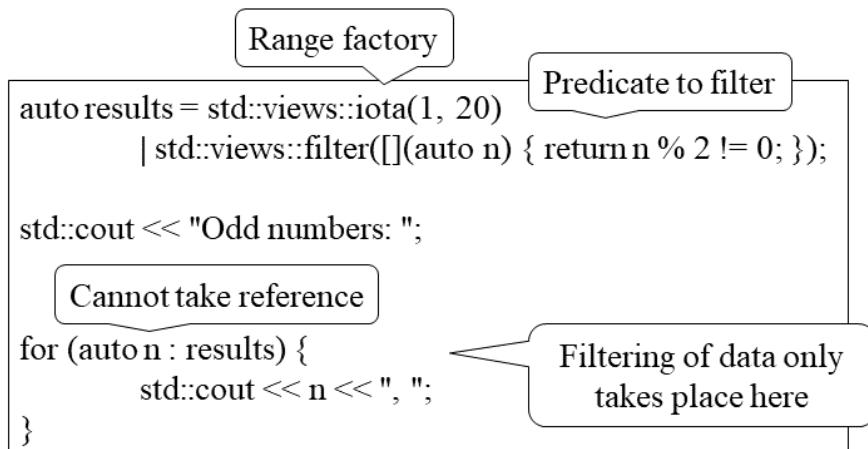
for (const auto& n : results) {
 std::cout << n << ", ";
}
```

Diagram annotations:

- A callout box labeled "Predicate used to filter" points to the `std::views::filter` call.
- A callout box labeled "Reference to value" points to the `n` in the loop variable declaration `const auto& n`.
- A callout box labeled "Filtering of data only takes place here" points to the `return n % 2 != 0;` condition in the predicate.

# Pipelining (lazy evaluation)

- Sequences of data can be processed by pipelining, using range adaptors:



## Common

- One of the issues mentioned earlier was that of algorithms requiring pairs of iterators of the same type (common range)
  - This may not be the case for some range adapters:

```
std::vector<int> data{ 1,2,3,4,5 };
std::vector<int> data2;
auto take_view = data
 | std::views::take_while([](auto n) {return n < 4; })
 | std::views::common; Make common_range
std::copy(take_view.begin(), take_view.end(), std::back_inserter(data2));
```

# Range Adaptors for Elements

- Where sequences of objects are processed within a pipeline, elements can be selected:

| Adaptor  | Description                         |
|----------|-------------------------------------|
| elements | Select specific element from tuples |
| keys     | Select key from pair                |
| values   | Select value from pair              |

```
std::map<int, std::string> data{ {1, "one"}, {2,"two"}, {3, "three"} };

auto keys_results = data | std::views::keys;

std::ostream_iterator<int> osi_int(std::cout, ", ");
std::ranges::copy(keys_results, osi_int); 1, 2, 3,
```

# Range Adaptor for Values

- Similarly values can be obtained from pairs:

```
std::map<int, std::string> data{ {1, "one"}, {2,"two"}, {3, "three"} };

auto values_results = data | std::views::values;

std::ostream_iterator<std::string> osi_string(std::cout, ", ");
std::ranges::copy(values_results, osi_string); one, two, three,
```

# Range Adaptor for Elements

- Sequence of tuples to be processed:

```
std::vector<std::tuple<int, std::string>> data{ {1, "one"}, {2,"two"}, {3, "three"} };

auto elements0_results = data | std::views::elements<0>;
auto elements1_results = data | std::views::elements<1>;

std::ostream_iterator<int> osi_int(std::cout, ", ");
std::ostream_iterator<std::string> osi_string(std::cout, ", ");

std::ranges::copy(elements0_results, osi_int);
std::cout << std::endl;
std::ranges::copy(elements1_results, osi_string);
std::cout << std::endl;
```

**1, 2, 3,**  
**one, two, three,**

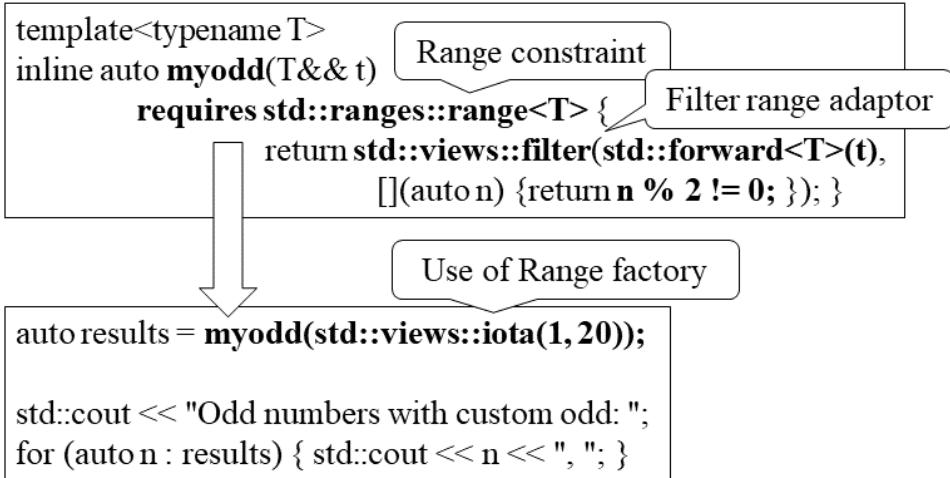
# Custom Range Adaptor

- The following is a step towards creating and range adaptor to filter out even numbers:

```
template<typename T>
inline auto myodd(T&& t) {
 requires std::ranges::range<T> {
 return std::views::filter(std::forward<T>(t),
 [](auto n) {return n % 2 != 0; });
 }
}

auto results = myodd(std::views::iota(1, 20));

std::cout << "Odd numbers with custom odd: ";
for (auto n : results) { std::cout << n << ", "; }
```



## Custom Range Adaptor - continued

- The following is a step towards creating and range adaptor to filter out even numbers:

```
template<typename T>
inline auto myodd2(T&& t) {
 requires std::ranges::range<T>
 { return std::forward<T>(t) | std::views::filter([](auto n) {return n % 2 != 0;}); }
}

auto results = myodd2(std::views::iota(1, 20));
std::cout << "Odd numbers with custom odd: ";
for (auto n : results) { std::cout << n << ", "; }
```

## Custom Range Adaptor - continued

- The following is a step towards creating and range adaptor to filter out even numbers:

```
constexpr inline auto myodd3()
{ return std::views::filter([](auto n) {return n % 2 != 0;}); }

auto results = std::views::iota(1, 20)
 | myodd3()

std::cout << "Odd numbers with custom odd: ";
for (auto n : results) { std::cout << n << ", "; }
```

# Threading

- This section considers the new thread related features:
  - jthread
  - stop\_source and stop\_token
  - counting\_semaphore
  - latch
  - barrier

# Thread class

- The ‘thread’ class provides the ability to create a single thread of execution
  - If ‘join’ is required, explicit call to the ‘join’ is required
  - All paths would require to ‘join’
    - Complicated by possible throwing of exceptions
  - The thread starts when the thread is created (subject to scheduling)

## jthread

- The ‘jthread’ class is similar to thread except provides for ‘RAII’
- ‘jthread’ joins on destruction
  - No need to add call to ‘join’ on multiple paths
- ‘jthread’ also supports cancellation
  - Thread callable can take a ‘stop\_token’

## jthread ‘joining’

- The following use of jthread will complete successfully:

```
int main()
{
{
 std::jthread th{ []() {
 for (size_t i = 0; i < 10; i++)
 {
 std::cout << i << std::endl;
 }
 } }; // No need to call 'join' member function
 std::cout << "Hello World!\n";
}
```

## stop\_source and stop\_token

- The stop\_token can be used to control the termination of a thread
  - A stop\_token must be associated with a stop\_source
  - When creating a thread the ‘callable’ can take a stop\_token
  - Request for stop can be made and checked for within thread

## stop\_token and stop\_callback

- The stop\_token allows ‘observation’:

| Member function | Description                                              |
|-----------------|----------------------------------------------------------|
| stop_requested  | Returns Boolean to indicate if a stop has been requested |
| stop_possible   | Returns Boolean to indicate if a stop is possible        |

- A stop ‘callback’ can be created to execute when a stop is requested:

Associate with stop\_token for thread

```
std::stop_callback sc(th.get_stop_token(),
[]() { std::cout << "Callback for stop!!" << std::endl; });
```

Lambda callback, called when stop is requested

## **stop\_source**

- A `stop_source` has an associated ‘stop state’
  - Use to indicate if threads should stop

| <b>Member function</b>      | <b>Description</b>                                                                    |
|-----------------------------|---------------------------------------------------------------------------------------|
| <code>request_stop</code>   | Request stop of thread and may change stop state (subsequent requests have no effect) |
| <code>get_token</code>      | Returns a <code>stop_token</code>                                                     |
| <code>stop_requested</code> | Returns Boolean to indicate if a stop has been requested                              |
| <code>stop_possible</code>  | Returns Boolean to indicate if a stop is possible                                     |

## **counting\_semaphore**

- Counting semaphore is a lightweight synchronisation object to protect resource(s)
  - Contains an internal count
  - Count can be ‘acquired’ (decremented)
  - Count can be ‘released’ (incremented)
- ‘`binary_semaphore`’ equivalent to `counting_semaphore<1>`
- Count can be configured at declaration and initialisation

# counting\_semaphore

- Member functions:

| Member function   | Description                                                           |
|-------------------|-----------------------------------------------------------------------|
| release           | Increment count, to allow another thread to acquire count             |
| acquire           | Decrement count, if available, otherwise blocks                       |
| try_acquire       | Attempts to acquire and returns Boolean indicating success or failure |
| try_acquire_for   | Attempts to acquire for a given duration                              |
| try_acquire_until | Attempts to acquire until a given time                                |

# std::latch

- The latch is a single use barrier
- The latch has a count which is decremented
  - Count initialised on creation
- Threads wait on the latch until the count is zero
  - Thus creating a form of barrier

# std::latch members

- The member functions are:

| Member functions | Description                                            |
|------------------|--------------------------------------------------------|
| count_down       | Decrement the counter                                  |
| try_wait         | Returns Boolean indicating count is zero               |
| wait             | Wait on count to become zero                           |
| arrive_and_wait  | Decrement the counter and wait on count to become zero |

- Constant:

| Constant | Description                        |
|----------|------------------------------------|
| max      | Implementation max value for count |

## latch example

- The following illustrates the use of a latch:

```
using std::chrono_literals::operator"" s;
{
 std::latch l(3); Latch with count of three
 std::jthread th1{ [&l](int n) {
 for (size_t i = 0; i < n; i++) {
 std::this_thread::sleep_for(1s);
 std::cout << i << std::endl;
 }
 std::cout << "Completed..." + std::to_string(n) << std::endl;
 l.count_down(); Decrement Total of three threads decrementing latch
 },20 };
 std::jthread th2{ [&l](int n) { ... } };
 l.wait(); Wait for count to reach zero
}
```

# std::barrier

- Barrier are effectively multi use, whereas latch is single use
  - Meaning that threads can go through multiple stages
  - The threads wait for all eligible threads to reach the barrier
  - A completion function is called when threads reach the barrier

## std::barrier members

- The member functions are:

| Member functions | Description                                                                                        |
|------------------|----------------------------------------------------------------------------------------------------|
| arrive           | Arrival of thread ‘at’ the barrier (decrement count)                                               |
| wait             | Blocks at sync point for phase completion                                                          |
| arrive_and_wait  | Arrival of thread ‘at’ the barrier (decrement count) and blocks at sync point for phase completion |
| arrive_and_drop  | Decrement expected count and expected count                                                        |

- Constant:

| Constant | Description                        |
|----------|------------------------------------|
| max      | Implementation max value for count |

# barrier example

- The following illustrates the use of a barrier:

```
std::array<std::string, 5> messages{ "one", "two", "three", "four", "five"};
Completed 'callable' noexcept required Data to be passed to thread
auto completed = []() noexcept {
 static std::string output= "Completed...\n";
 std::cout << output << std::endl;
 output= "Done...";
};
barrier Signed size of collection Completed 'callable'
std::barrier sync(std::ssize(messages), completed);
```

# barrier example - continued

- The following illustrates the use of a barrier:

```
Barrier
auto work = [&sync](std::string msg) {
 std::cout << "Working on " + msg + "\n";
 sync.arrive_and_wait();
 std::cout << "Continuing work on " + msg + "\n";
 sync.arrive_and_wait();
};

std::vector<std::jthread> threads;
for (size_t i = 0; i < std::size(messages); i++)
{
 threads.push_back(std::jthread(work, messages[i]));
}
```

## barrier - output

- The following would be the resulting output from the barrier example:

**Working on one  
Working on two  
Working on four  
Working on five  
Working on three  
Completed...**

**Continuing work on three  
Continuing work on four  
Continuing work on five  
Continuing work on two  
Continuing work on one  
Done...**

## Atomic

- Many new features have been added in relation to Atomic:
  - Atomic shared\_ptr
  - Atomic – Wait and Notify
  - Atomic reference

## Atomic shared\_ptr

- The management block for std::shared\_ptr was thread safe, however the use of the pointer was not thread safe
  - Previously synchronisation could be used (std::mutex) to protect the pointer
  - std::atomic\_load and std::atomic\_store now deprecated (for std::shared\_ptr)
- Can now use std::atomic<std::shared\_ptr<>> with member functions (specialised)

## Atomic std::shared\_ptr

- Shared pointers can now be atomic
  - Atomic shared\_ptr cannot be copy or move constructable or assignable

```
{
 std::atomic<std::shared_ptr<int>> ptr{ std::make_shared<int>(7) };
 std::cout << "Current: " << *ptr.load() << std::endl;

 std::shared_ptr<int> val{std::make_shared<int>(8)};
 ptr.store(val);

 auto result = ptr.load();

 std::cout << "Replaced: " << *result << std::endl;
}
```

# Atomic Wait and Notify

- Atomic variables now support the member functions:

| Member function | Description                                                                                                                        |
|-----------------|------------------------------------------------------------------------------------------------------------------------------------|
| wait            | Waits on an atomic variable until notified<br>Wait starts if the ‘old’ value matches the value currently stored, otherwise returns |
| notify_one      | Notifies and ends the wait of one waiting thread                                                                                   |
| notify_all      | Notifies and ends the wait of all waiting thread                                                                                   |

# Producer Thread

- The following illustrates the use of atomic wait:

```
std::atomic_int val{ 0 };
std::jthread producer{ [&val]() {
 int cnt{0};

 while (cnt < 10) {
 val.store(cnt++);
 val.notify_all();
 auto second = std::chrono::seconds{ 1 };
 std::this_thread::sleep_for(second);
 }
};
```

```
std::jthread consumer{ [&val]() {
 int cnt{0};
 while (true) {
 val.wait(cnt);
 cnt=val.load();
 std::cout << cnt << std::endl;
 }
};
```

# Atomic Reference

- The following illustrates the use of atomic wait:

```
int val{ 0 };
std::atomic_ref ref{ val };
std::jthread producer{ [&ref]() {
 int cnt{ 0 };

 while (cnt < 10) {
 ref = cnt++;
 ref.notify_all();
 auto second = std::chrono::seconds{ 1 };
 std::this_thread::sleep_for(second);
 }
} };
```

```
std::jthread consumer{ [&ref]() {
 int cnt{ 0 };
 while (true) {
 ref.wait(cnt);
 cnt = ref;
 std::cout << cnt << std::endl;
 }
} };
```

# General C++ 20 Features

- C++ 20 has introduced many new features:
  - Using enum
  - Lambda Capture of ‘this’
  - constexpr and consteval
  - Constinit
  - Designated Initialiser
  - const virtual functions
  - Conditional explicit
  - [[likely]] and [[unlikely]] attributes

# Using enum

- It is now possible to introduce an enum into scope with a using statement:

```
enum class Number { One = 1, Two = 2, Three = 3 };

struct SomeNumbers
{
 using enum Number; Number introduced
}; into structure

int main()
{
 SomeNumbers sn; Access enum value
 in a number of ways

 std::cout << (int)sn.One << std::endl;
 std::cout << (int)SomeNumbers::Two << std::endl;
}
```

# Using enum - continued

- The using for enum is useful within switch statements:

```
enum class Number { One = 1, Two = 2, Three = 3 };
int main(){
 Number num{ Number::Two };
 switch (num){
 case Number::One: →
 std::cout << "One" << std::endl;
 break;
 case Number::Two:
 std::cout << "Two" << std::endl;
 break;
 ...
 }
}

switch (num){
 using enum Number;
 case One:
 std::cout << "One" << std::endl;
 break;
 case Two:
 std::cout << "Two" << std::endl;
 break;
 ...
}
```

# Lambda Capture of ‘this’

- Whilst this change is in one sense minor, it could potentially be a breaking change:
  - Prior to C++ 20 the use of ‘[=]’ could capture everything from the scope, including ‘this’ for member functions
    - Thus it is proposed to deprecate the use of [=] for the capture of ‘this’
  - In C++ 20 to produce the same effect would require the use of ‘[=,this]’

# constexpr and consteval

- constexpr was introduced with C++11 to allow evaluation at compile time of expressions involving function calls
  - C++11 placed a number of restrictions on the definition of the function
  - These restrictions were relaxed in later standards to allow use of assignment and loops
  - constexpr functions may also be used at runtime where an argument is a variable

# Return Type

- The return type and parameters to `constexpr` functions have the named requirement `LiteralType`:
- Requirements include being one of:
  - `void`
  - Scalar or reference type
  - Array of literal type
  - Class (with various requirements)!

## `constexpr` and `consteval` continued...

- `constexpr` was technically required on both the function and the variable being initialised
  - Although some compilers were a little relaxed
- `constexpr` is used widely within the standard library
  - Many algorithms are now `constexpr`
- `consteval` is similar to `constexpr` but with the restriction that a `consteval` function can only be evaluated at compile time (immediate function)
- An ‘immediate function’ is one which is evaluated at compile time to return a constant

## constexpr in C++ 20

- The use of constexpr has now been expanded further
  - virtual function can now be constexpr
  - Whilst constexpr functions require to be self contained, they can now use:
    - Dynamic memory allocation
      - Using only global ‘new’ and ‘delete’
    - try/catch blocks (although be noexcept)
    - ‘dynamic\_cast’ and ‘typeid’

## constexpr virtual functions

- Clearly there will be limitations on how this can be applied!

```
struct A {
 constexpr virtual int value() const { return 7; }
};
struct B : A {
 constexpr int value() const override { return 9; }
};
constexpr int add(const A& x, const A& y) {
 return x.value() + y.value();
}

int main()
{
 constexpr auto result = add(B{}, A{});
 ...
}
```

## std::string and std::vector

- The changes in support for constexpr have now made it possible to use both std::string and std::vector in constexpr functions
  - Now possible to use within constexpr functions
  - The following can now be treated as constexpr:

```
constexpr int sum_vec(const std::vector<int>& data)
{
 int result{};
 for (const auto& n : data) { result += n; }
 return result;
}
```

```
constexpr auto result = sum_vec(std::vector{1, 2, 3, 4});
```

## constexpr Algorithms

- Throughout the C++20 standard library constexpr is now widely used
- Many Algorithm are now constexpr

```
constexpr int sum(const std::vector<int>& data){
 return std::reduce(data.cbegin(), data.cend(), 0,
 [](auto a, auto b){return a+b;});
}
int main()
{
 constexpr int result = sum(std::vector{3,4,5,6});
 ...
}
```

## constinit

- The variable has ‘static’ initialisation
- Whilst appearing similar constinit is used for initialising ‘static’ and thread local values at compile time
- The keywords constexpr, consteval and constinit cannot be used in combination

## constinit

- Initialising a static variable:

```
constexpr int factorial(int n)
 { return n==0? 1 : n*factorial(n-1);}

int main()
{
 int n = 2;

 static constinit int fact7 = factorial(7);

 // static constinit int factn = factorial(n);

 return 0;
}
```

Variable parameter mean is cannot  
be evaluate at compile time

# Designated Initialiser

- For simple aggregate types (struct) designated initialisers allow initialisation of members by member name:

```
struct some_data {
 int a{1}, b{2};
};
int main()
{
 some_data sd{ .b = 42 };
 ...
}
```

Named member to be given value.  
Must be named in declared order

Value of ‘a’ is default

## explicit

- Constructors and type conversion operators can be used for implicit type promotion or type conversion
- The compiler is allowed to use one level of these in statements and expressions
  - This can be useful to avoid cluttering code with constructor calls and casts
- However, especially where operators are defined, ambiguities can arise

# explicit

- The explicit keyword has been available for use on one parameter constructors for a long time
  - Which forces an explicit call to constructor
- The explicit keyword was made available for type conversion operators from C++ 11
  - Which forces an explicit cast for type conversion
- The application of the explicit keyword can be used to resolve ambiguities mentioned

## Conditional explicit

- Explicit can now be conditional, which means that within template classes, one parameter constructors and type conversion operators can have finer grain control:

```
template<typename T>
class Num {
 ...
public:
 ...
 explicit(std::is_integral_v<T> || std::is_floating_point_v<T>)
 operator T() const {
 return ... }
 };
}
```

Use appropriate  
type traits

## [[likely]] and [[unlikely]] attributes

- The [[likely]] and [[unlikely]] attributes are to be used to give a hint to the compiler for optimisation purposes
  - Can be applied within conditional statements to indicate likely and unlikely ‘branches’

```
int num{};
std::cin >> num; Hint that branch is unlikely!
if(num > 0)[[unlikely]]{
 std::cout << "Greater than zero!" << std::endl;}
else{
 std::cout << "Less than or equal zero!" << std::endl;
}
```

## Switch with Attributes

- [[likely]] and [[unlikely]] can be used within a switch statement:

```
int num{};
std::cin >> num;
switch(num){
 [[unlikely]]case 1:std::cout << "One" << std::endl;
 break;
 [[likely]]case 2:std::cout << "Two" << std::endl;
 break;
 case 3:std::cout << "Three" << std::endl;
 break;
}
```

- Cannot use [[likely]] and [[unlikely]] on same statement

## General C++ 20 Features - continued

- C++ 20 has introduced many new features:
  - Spaceship Operator
  - `bind_front` function
  - Source Location
  - Initialisation in Range based For
  - Shared pointer for arrays
  - Removing Items
  - Format Library

## Spaceship Operator

- This section is to introduction to the new spaceship operator:
  - Spaceship Operator  $\leqslant \geqslant$
  - Relational Operators
  - Expression Rewriting
  - Simple Class with Spaceship
  - Operator Symmetry
  - Custom Spaceship Operator
  - Spaceship Operator Functions

# Spaceship Operator <=>

- The spaceship operator (or three way comparison) is new to C++20 and reduces the need to explicitly provide relational operators
  - Traditionally many operators could be provided:

| Operator   | Description           |
|------------|-----------------------|
| operator<  | Greater than          |
| operator<= | Greater than or equal |
| operator>  | Less than             |
| operator>= | Less than or equal    |
| operator== | Equals                |
| operator!= | Not equals            |

# Relational Operators

- The traditional way relational operators are defined is as non member functions (in order to be symmetric)
  - Thus defined as friend operator function or global function outside of the class
  - Define one or two and define others in terms of those defined (in order to avoid mistakes)!
  - operator< should satisfy ‘strict weak ordering’ in order to be used with containers and algorithms

# Expression Rewriting

- Defining the Spaceship operator now uses another new feature of C++20
- If the expression ‘`a < b`’ is written for variables of type with the spaceship operator, this expression is rewritten as ‘`(a <= b) < 0`’
- The equality and inequality are defined separately but nevertheless generated when the spaceship operator is provided

## Simple class with Spaceship

- Defined class:

```
class ValueClass
{
 int _val{};
public:
 ValueClass(int val):_val{val} {}
 auto operator<=>(const ValueClass& rhs)const = default;
};
```

- Generated:

```
class ValueClass
{
 int _val;
public:
 inline ValueClass(int val) : _val{val} {}
 inline constexpr std::strong_ordering operator<=>(const ValueClass & rhs)
 const noexcept = default;
 inline constexpr bool operator==(const ValueClass & rhs) const = default;
};
```

Strong ordering as  
comparing integral types

## <=> Generated Code

- In the previous example the return type of the operator is std::strong\_ordering
  - This was due to the one data member being int
- If the data members include a floating point number the return type will be std::partial\_ordering as some values are not comparable
- Where a class has multiple data members the operator compares the data members in order

## Class Containing Double

- Use the default implementation of <=>:

```
struct TheDouble {
 double d;
 auto operator<=>(const TheDouble& f) const = default;
};
```

```
int main() {
 TheDouble td1{ 34.2 }, td2{ NAN };

 auto result = td1 <=> td2;
```

NAN or INFINITE could be tried

‘unordered’ for NAN and ‘greater’ for INFINITE

# Operator Symmetry

- When defining operators manually symmetry had to be considered
- Expression rewriting takes this into account,  
e.g.

```
ValueClass vc1{1},vc2{2};

if(1 < vc2){ std::cout << "1 less than 2!\n" << std::endl; }
```

- Generated code:

```
ValueClass vc1 = ValueClass{1};
ValueClass vc2 = ValueClass{2};

if(operator<(__cmp_cat::__unspec(0), vc2.operator<=>(ValueClass(1)))) {
 std::operator<<(std::cout, "1 less than 2!\n").operator<<(std::endl);
}
```

Right hand  
operand on left

Type  
promotion used

## std::strong\_ordering

- The previous generated code return type for the spaceship operator is not Boolean!
  - Return type was std::strong\_ordering
- std::strong\_ordering has four static constants

| Constants  | Description                                                    |
|------------|----------------------------------------------------------------|
| less       | Represents a less-than relationship (ordered before)           |
| greater    | Represents a greater-than relationship (ordered after)         |
| equal      | Representing equivalence (not ordered before or ordered after) |
| equivalent | Representing equivalence – same as equal                       |

# Custom Spaceship Operator

- It is possible to define a custom operator `<=>`:

```
class ValueClass
{
 int _val{};
public:
 ValueClass(int val):_val{val} {}
 constexpr std::strong_ordering operator<=>(const ValueClass& rhs) const {
 if(_val < rhs._val) return std::strong_ordering::less;
 else if(_val > rhs._val) return std::strong_ordering::greater;
 else return std::strong_ordering::equal;
 }
 constexpr bool operator==(const ValueClass& rhs) const {
 return _val == rhs._val;
 }
};
```

**operator!= will  
be generated**

# Orderings

- All orderings should be:

|             | Description                                         |
|-------------|-----------------------------------------------------|
| irreflexive | $a < a$ is false                                    |
| asymmetric  | $a < b$ true, implies $b < a$ false                 |
| transitive  | $a < b$ true and $b < c$ true, implies $a < c$ true |

- A number of orderings are defined:

| Ordering         | Constants used                                         |
|------------------|--------------------------------------------------------|
| Strong Ordering  | greater/equal/equivalent/less (substitutable if equal) |
| Weak Ordering    | greater/equivalent/less (not substitutable)            |
| Partial Ordering | greater/equivalent/less/unordered                      |

Equivalent means neither ordered before or after!

Incomparable value!

# Spaceship Operator Functions

- New functions have been added which take the result from a spaceship operator

| Function | Description           |
|----------|-----------------------|
| is_gt    | Greater than          |
| is_gteq  | Greater than or equal |
| is_lt    | Less than             |
| is_lteq  | Less than or equal    |
| is_eq    | Equals                |
| is_neq   | Not equals            |

## Spaceship operator Example

- Alternative ways spaceship operator could be used explicitly:

```
struct some_data {
 int a{1};
 int b{2};
 auto operator<=>(const some_data&)const= default;
};

int main(){
 some_data sd;
 std::strong_ordering result = some_data{2, 3 } <=> sd;
 if (result == std::strong_ordering::greater)
 std::cout << "greater than..." << std::endl;
 if (std::is_gt(result)) std::cout << "greater than..." << std::endl;
}
```

## `std::bind_front`

- The `std::bind` function provides a very flexible means of creating a new ‘callable’ from an existing ‘callable’
- `std::bind_front` provides a more limited, but useful simple means of creating a new ‘callable’, by binding just the first arguments
  - The remaining arguments can then be passed to the returned ‘callable’

## `bind_front` Example

- Using a simple function:

```
int calc(int a, int b, std::string c, std::string d) {
 return a + b - std::stoi(c) * std::stoi(d);
}
int main(){
 auto part = std::bind_front(calc, 3, 7);
 std::cout << "Result: " << part("2", "3") << std::endl;;
}
```

Bind first two arguments

Pass remaining arguments

## bind\_front Member Function Example

- Using a simple function:

```
class DataClass {
 int _a;
public:
 DataClass(int a) :_a{ a } {}
 int calc(int b, int c) const { return _a - b + c; }
};
int main(){
 DataClass dc{ 10 };
 auto member = std::bind_front(&DataClass::calc, dc);
 std::cout << "Result: " << member(1, 2) << std::endl;
}
```

The code shows a simple class DataClass with a constructor and a member function calc(). In the main() function, a bind\_front operation is performed on the calc() method, binding it to the object dc. The resulting member variable is then called with arguments 1 and 2, outputting the result to the console.

## std::source\_location

- std::source\_location allows obtaining location within code:

| Member function | Description                                                                      |
|-----------------|----------------------------------------------------------------------------------|
| current         | Returns source_location corresponding to where the call to 'current' takes place |
| line            | Returns line number where source location obtained                               |
| column          | Returns column number where source location obtained                             |
| file_name       | Returns file name where source location obtained                                 |
| function_name   | Returns function name where source location obtained                             |

# Source Location - Example

- Function to obtain location:

```
void logit(const std::string& msg,
 std::source_location sl = std::source_location::current()) {

 std::cout << msg << std::endl;
 std::cout << "Function: " << sl.function_name()
 << ", Line number: " << sl.line() << std::endl;
}
```

Called when function is called

- Using the default parameter allows obtaining position for ‘logit’ function call

# Initialisation in Range based For

- C++ 17 introduced initialisation within both switch and while
  - Now also available within range based for:

```
auto get_data() {
 return std::vector<int>{7, 6, 5, 4, 3, 2, 1};
}
int main()
{
 for (auto data = get_data(); const auto & n: data) {
 std::cout << n << std::endl;
 }
}
```

## std::shared\_ptr for Arrays

- Shared pointers to arrays can now be created using `make_shared`, e.g.:

```
class A {
public:
 ~A() { std::cout << "Destroying A..." << std::endl; }
};
```

```
int main()
{
 auto ptr = std::make_shared<A[]>(5);

}
```

Output:+

```
Destroying A...
Destroying A...
Destroying A...
Destroying A...
Destroying A...
```

## Removing Items (pre C++ 20)

- Removing items from a container creates an interesting challenge, as it is important not to invalidate iterators
- There are a number of options for removing items:
  - Removing by moving values up, this does not change the size of the container
  - Remove copy, to populates a new container with only wanted items

## std::remove and std::remove\_if

- std::remove and std::remove\_if simply moves values up in the container
  - The container size remains the same
  - Returns new ‘end’ iterator
    - Iterator to value after last wanted item.

```
void remove_value(std::vector<int>& data, int value)
{
 auto end = std::remove(std::begin(data), std::end(data), value);

 auto newEnd = std::remove_if(std::begin(data), std::end(data),
 std::bind(std::less<int>(), std::placeholders::_1, 3));
}
```

Compared using  
== operator

## C++ 20 erase and erase\_if

- For std::vector the erase member function could then be used
- With C++ 20 there is std::erase and std::erase\_if algorithms for vector and string

```
std::cout << "The C++ 20 way:" << std::endl;
std::vector data{ 1,2,3,4,5,6,7,8,1,2,3,4,5,6,7,8 };

std::erase(data, 5);

std::ostream_iterator<int> osi(std::cout, ", ");
std::ranges::copy(data, osi);
```

The C++ 20 way:  
1, 2, 3, 4, 6, 7, 8, 1, 2, 3, 4, 6, 7, 8,

# Format Library

- This section introduces the format library:
  - Format Library Introduction
  - std::format
  - std::vformat
  - Formatting String
  - Numbered Placeholders
  - Format String Specification
  - Format Integral Types
  - Format Floating Point Types

## Formatting Library Introduction

- C++ already has the stream library and the legacy printf function, however the new formatting library provides an extensible alternative
  - Complements stream library and reuses some features
  - std::format allows formatting of a std::string using a format string and arguments

# Formatting Library continued

- The library contains many functions and overloads for formatting strings
  - Overloads for wide characters and locale
- String can be created in a number of ways:

| Function   | Description                                                                          |
|------------|--------------------------------------------------------------------------------------|
| format     | Returns a string using format and arguments                                          |
| format_to  | Outputs a string using an iterator                                                   |
| vformat    | Returns a string using format (view) and arguments (using make)!                     |
| vformat_to | Outputs a string using an iterator, from a format (view) and arguments (using make)! |

## std::format

- std::format can be used in a number different ways (include <format>):

```
auto n = 2;
auto d = 3.2;
auto f = 65.3f;
```

string                  Format string                  Arguments to populate {}

auto message = std::format( "int: {}, double: {}, float: {}", n, d, f);

{} placeholders filled in order

# std::vformat

- std::vformat can be used in a number different ways (include <format>):

```
auto n = 2;
auto d = 3.2;
auto f = 65.3f;

string Format string
std::string fint = "int: {}, double: {}, float: {}";
 string_view passed

auto message = std::vformat(fint, std::make_format_args(n, d, f));
 Make format args store

std::cout << message << std::endl;
```

# Formatting String

- Formating strings is straightforward:

```
using std::string_literals::operator "" s;
auto b{ true };
auto ch {'g'};
auto str{ "Greetings"s};
try {
 auto message = std::format("s:{0:s},s:{1:c},s:{2:s}", b, ch, str);
 std::cout << message << std::endl;
}
catch (const std::format_error& fe) {
 std::cout << fe.what() << std::endl;
}
```

# Numbered Placeholders

- Placeholders can be numbered
- The format function can throw an exception

```
auto n = 2;
auto d = 3.2;
auto f = 65.3f;
try {
 auto message = std::format("float: {1}, int: {0}, double: {2}", n, d, f);
 std::cout << message << std::endl;
}
catch (const std::format_error & fe) {
 std::cout << fe.what() << std::endl;
}
```

Format string      Arguments to populate {}, additional arguments will be ignored

Number to identify argument

# Format String Specification

- The specification has many options ([] indicates optional):

[fill and align][sign][#][0][width][precision][L][type]

– Fill and align

| Symbol | Meaning                               |
|--------|---------------------------------------|
| <      | Left justified (default none numeric) |
| >      | Right justified (default numeric)     |
| ^      | Centered                              |

– Sign

| Symbol | Meaning                               |
|--------|---------------------------------------|
| +      | Both negative and positive sign shown |
| -      | Only negative sign shown (default)    |
| space  | Space for positive sign               |

# Format String Specification continued

- The specification has many options ([] indicates optional):

[fill and align][sign][#][0][width][precision][L][type]

- [#] indicates prefix should be shown for Boolean, Octal and Hexadecimal
- [0] indicates pad with zeros
- [width] is the minimum field width
- [precision] is a ‘.’ and the number of digits
- [L] formatting specific to locale

## Format Specifiers

- Following are some of the format specifiers:

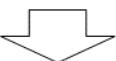
| Format specifier | Type                 | Description (# to show prefix)    |
|------------------|----------------------|-----------------------------------|
| none, s          | string               | String copied                     |
| b                | Integral (not chars) | Binary format                     |
| B                | “                    | Binary format, but prefix 0B      |
| c                | “                    | Character                         |
| d                | “                    | Decimal format                    |
| o                | “                    | Octal format                      |
| x                | “                    | Hexadecimal format                |
| X                | “                    | Hexadecimal format, but prefix 0X |
| none             |                      | Defaults to ‘d’                   |

# Format Integral Type

- Use of format characters:

```
auto n = 287;

try {
 auto message = std::format(
 "b:{0:#b},B:{0:#B},d:{0:d},o:{0:#o},x:{0:#x},X:{0:#X}", n);
 std::cout << message << std::endl;
}
catch (const std::format_error& fe) {
 std::cout << fe.what() << std::endl;
}
```



**b:0b100011111, B:0B100011111, d:287, o:0437, x:0x11f, X:0X11F**

# Format Specifiers continued

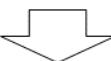
- Following are some of the format specifiers:

| Format specifier | Type           | Description                             |
|------------------|----------------|-----------------------------------------|
| a                | Floating point | Floating point number                   |
| A                | “              | Floating point number (capital letters) |
| e                | “              | Scientific notation                     |
| E                | “              | Scientific notation, Uses E             |
| f                | “              | Floating point number                   |
| F                | “              | Floating point number (capital letters) |
| g                | “              | General number format                   |
| G                | “              | General number format (capital letters) |
| none             | “              | Similar to g                            |

# Format Floating Point Type

- Use of format characters:

```
auto d = 2.2;

try {
 auto message = std::format(
 "a:{0:a},A:{0:A},e:{0:e},E:{0:E},f:{0:f},F:{0:F}", d);
 
 a:1.199999999999ap+1,A:1.199999999999AP+1,
 e:2.200000e+00,E:2.200000E+00,f:2.200000,F:2.200000
 std::cout << message << std::endl;
}
catch (const std::format_error& fe) {
 std::cout << fe.what() << std::endl;
}
```

# Specific Floating Point

- Use of options:

```
auto d{ 234.235 };

try {
 auto message = std::format("double precision number: '{0:^+12.4f}'", d);
 
 double precision number: '+234.2350 '
 std::cout << message << std::endl;
}
catch (const std::format_error& fe) {
 std::cout << fe.what() << std::endl;
}
```

# Print Library (C++23)

- This section gives an overview of the new C++23 print library:
  - Introduction to ‘print’
  - Printing with `<iostream>`
  - Printing with `<print>`

## Introduction to ‘print’

- The format library provides a means of creating a string from a number of arguments and a format string
- This formatted string could then be output using streams
- The new ‘print’ capabilities provide functions which can be called to output arguments to an output stream in a formatted way

# Printing with <iostream>

- The header file <iostream> now provides a number of new functions:

| Function | Description                                           |
|----------|-------------------------------------------------------|
| print    | Outputs to an ostream a number of formatted arguments |
| println  | Like print, but appends a new line character          |

```
{
int n{5}, square{n * n};

std::println(std::cout, "The square of {0} is {1}", n, square);
}
```

**The square of 5 is 25.**

# Printing with <print>

- The new header file <print> provides:

| Function | Description                                       |
|----------|---------------------------------------------------|
| print    | Outputs to stdout a number of formatted arguments |
| println  | Like print, but appends a new line character      |

```
{
int n{5}, square{n * n};

std::println(std::cout, "The square of {0} is {1}", n, square);
}
```

**The square of 5 is 25.**

## Printing with <print>

- Print and println are overloaded to take file:

```
{
FILE* aFile = new FILE();

fopen_s(&aFile, "C:/Temp/Info.txt", "w");
if (aFile) {
 int n{ 5 }, square{ n * n };

 std::println(aFile, "The square of {0} is {1}.", n, square);
 std::fclose(aFile);
}
}
```

## C++ 23 Operators

- This section gives an introduction to some of the new features in C++ 23:
  - Static Operators
  - Usage of static operator()
  - Static Lambdas
  - operator[]
  - operator[,]

# Static Operators

- C++ 23 now allows the operator() and operator[] (for consistency) to be declared static
  - The use of operator() to allow definition of function object has been available for a long time
  - However, where the type contains no members there seems little necessity for an instance function
  - Thus, it is now possible to define a the operator() as static
  - A compiler can then optimise away the pointer to ‘this’.

```
struct Square{
 static int operator()(int n){ return n*n; }
}
```

## Usage of static operator()

- Whilst the operator() may be static its usage is the same as for the instance operator:

```
{
 std::vector<int> data(10);
 std::ranges::iota(data, 1);

 std::ranges::transform(data, data.begin(), Square{});

 std::ranges::copy(data, std::ostream_iterator<int>{std::cout, ", "});
}
```



1, 4, 9 16, 24, 36, 49, 64, 81, 100,

# Static Lambdas

- A consequence of the static operator() is that the generated operator() for a lambda can now be static:
  - This cannot be done implicitly as this may break some existing code
  - Possible to opt-in to static operator by defining static lambda for capture-less lambdas

```
{ Empty capture
 auto fun = [](int n) static { return n*n; }
 std::cout << fun(12) << std::endl;
}
```

## operator[]

- For consistency it is also possible to define the operator[] as static for types with no data members:

```
struct Echo{
 static int operator[](int n){ return n; }
};
```

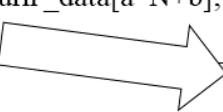
```
{
 Echo echo;

 std::cout << echo[7] << std::endl;
}
```

# operator[ , ]

- The operator[] can now support multiple comma separated ‘indices’:

```
template<unsigned N>
struct MyData {
 int *_data;
public:
 MyData(int *data):_data{data} {}
 ~MyData() { delete [] _data; }
 int operator[](int a, int b) { return _data[a*N+b]; }
};
```



```
{
 int *data = new int[] {1,2,3,4,5,6};
 MyData<3> md{data};

 for (size_t i = 0; i < 2; i++)
 {
 for (size_t j = 0; j < 3; j++)
 {
 std::cout << md[i,j] << ", ";
 }
 std::cout << std::endl;
 }
}
```

# C++ 23 Fold Algorithms

- This section gives an introduction to Fold Algorithms:
  - Fold Algorithms Introduction
  - Fold Expression (C++ 17)
  - Fold Algorithms
  - Fold Algorithms Example
  - Left Fold

# Fold Algorithms Introduction

- C++ 23 now provides algorithms performing fold operations
  - Fold expressions have previously been available by the use of variadic templates
  - Algorithms correspond to these fold expressions
  - Whereas the use of variadic templates require parameters to be passed, the algorithms allow the use of containers

# Fold Expressions (C++ 17)

- Fold expressions are used in conjunction with variadic templates
  - Makes implementing functions simpler
  - Does not require recursive implementation
- Four options:

| Fold Expression           | Description                                          |
|---------------------------|------------------------------------------------------|
| (pack op ...)             | Right unary expression                               |
| (... op pack)             | Left unary expression                                |
| (pack op ... op initexpr) | Right binary expression (operators must be the same) |
| (initexpr op ... op pack) | Left binary expression(operators must be the same)   |

# Fold Algorithms

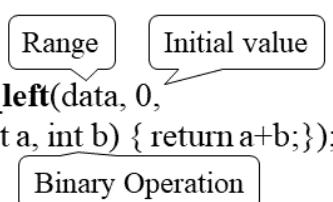
- Available algorithms are constrained algorithms and either take a range or iterator-sentinel pair:

| Algorithm                 | Description                    | Returns                 |
|---------------------------|--------------------------------|-------------------------|
| fold_right                | Takes initial value            | Result of fold          |
| fold_right_last           | Last element as initial value  | Result of fold          |
| fold_left                 | Take initial value             | Result of fold          |
| fold_left_first           | First element as initial value | Result of fold          |
| fold_left_with_iter       | Take initial value             | Result and end iterator |
| fold_left_first_with_iter | First element as initial value | Result and end iterator |

## Fold Algorithm Example

- The following illustrates the use of the `left_fold`:

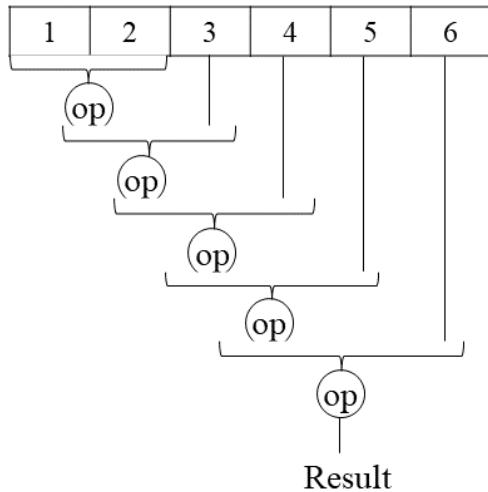
```
{\n std::vector<int> data(7);\n std::ranges::iota(data, 1);\n\n auto result = std::ranges::fold_left(data, 0,\n [](\n int a, int b) {\n return a+b;\n });\n\n std::cout << result << std::endl;\n}
```



# Left Fold

- The left fold:

- fold\_left:



# C++ Monad

- This section gives an introduction to implementing Monads in C++:
  - Monad Pattern
  - Implementing the Monad
  - Monad Usage
  - C++ 23 std::optional monadic functions
  - Monad in C++ 23
  - C++ 23 std::expected
  - Monad in C++ 23 (std::expected)

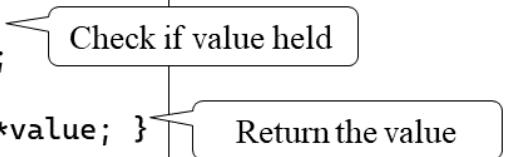
# Monad Pattern

- The monad pattern allows involves a ‘value’ and associated ‘computation’
- The use of monads allows data to be handled in a structured way
  - Monads can be used in a sequence in order to perform multiple computations
- Monads are typically use in functional programming and are particularly common in languages like Haskell

## Implementing the Monad

- Prior to C++ 23 monads could be implemented by wrapping the std::optional type:

```
template<typename T>
class Maybe {
private:
 std::optional<T> value;
public:
 Maybe() : value(std::nullopt) {}
 Maybe(T val) : value(val) {}
 bool has_value() const {
 return value.has_value();
 }
 T unwrap() const { return *value; }
 ...
};
```



## Implementing the Monad - continued

- An important member of the type is ‘bind’, which allow the computation if value exists:

```
template<typename T> class Maybe {
public:
 ...
 template <typename U>
 auto bind(std::function<Maybe<U>(T)> f) {
 if (value) { return f(*value); }
 else { return Maybe<T>(); }
 }
 static Maybe<T> return_value(T val) {
 return Maybe<T>(val);
 }
};
```

```
graph TD
 A[Execute the function] --> B[Return nullopt]
```

## Monad Usage - computations

- Given the possible computations defined below, these can be performed in sequence:

```
Maybe<int> add_five(int x) {
 return Maybe<int>::return_value(x + 5);
}

Maybe<int> divide_by_two(int x) {
 if (x % 2 == 0) {
 return Maybe<int>::return_value(x / 2);
 }
 else {
 return Maybe<int>();
 }
}
```

# Monad Usage

- Maybe can be used as:

```
int main() {
 Maybe<int> value = Maybe<int>::return_value(10);

 Maybe<int> result = value
 .bind<int>(add_five)
 .bind<int>(divide_by_two); Perform multiple computations

 if (result.has_value()) {
 std::cout << "Result: " << result.unwrap() << std::endl;
 }
 else {
 std::cout << "Failed." << std::endl;
 }
}
```

## C++ 23 std::optional monadic functions

- As of C++ 23 it is now possible to perform similar functionality using the new optional member functions:

| Member function | Description                                                                               |
|-----------------|-------------------------------------------------------------------------------------------|
| and_then        | Compute function if value exists and return optional otherwise return empty optional      |
| transform       | Return transformed value if it exists and return optional otherwise return empty optional |
| or_else         | Return optional if value exists otherwise result of give function                         |

# Monad is C++ 23

- The follow could now be used instead of the Maybe type
- First of all computations:

```
std::optional<int> add_five(int x) {
 return x + 5;
}
std::optional<int> divide_by_two(int x) {
 if (x % 2 == 0) {
 return x / 2;
 }
 else {
 return std::nullopt;
 }
}
```

# Monad in C++ 23 - continued

- Usage of the ‘monad’, implemented using optional:

```
int main() {
 std::optional<int> value = 10;
 auto result = value
 .and_then(add_five)
 .and_then(divide_by_two); Perform multiple computations

 if (result.has_value()) {
 std::cout << "Result: " << *result << std::endl;
 }
 else {
 std::cout << "Computation failed." << std::endl;
 }
}
```

## C++ 23 std::expected

- A variation on the monad implementation can be achieved using std::expected
- std::expected either holds a value or an error
  - Similar ‘monadic’ members are available to those on std::optional
  - On ‘failure’ the std::expended holds an error and not ‘empty’

## Monad is C++ 23 (std::expected)

- The follow could now be used instead of the Maybe type
- First of all computations:

```
std::expected<int, std::string> add_five(int x) {
 return x + 5;
}
std::expected<int, std::string> divide_by_two(int x) {
 if (x % 2 == 0) {
 return x / 2;
 }
 else {
 return std::unexpected(
 "Division by an odd number is not allowed");
 }
}
```

# Monad in C++ 23 - continued

- Usage of the ‘monad’, implemented using optional:

```
int main() {
 std::expected<int, std::string> initial_value = 10;

 auto result = initial_value
 .and_then(add_five)
 .and_then(divide_by_two); Perform multiple computations

 if (result) {
 std::cout << "Result: " << *result << std::endl;
 }
 else {
 std::cout << "Computation failed: "
 << result.error() << std::endl;
 }
}
```

## The End

Congratulations on completing this  
C++ course

# Appendix 1

## Modules

## Modules

- This section introduces the Modules feature:
  - Modules Introduction
  - Module Interface
  - Visual Studio Support
  - Module Interface Unit
  - Module Implementation Unit
  - Module Implementation Unit and Private
  - Module Partitions

# Modules Introduction

- Modules essentially replace header files
- Header files suffered from many problems
  - No encapsulation
  - Leakage of defines/macros
  - Inconsistency
  - Multiple builds – may result in high build times
  - Order significant
- Modules resolve all of these problems
  - Module interface - specific ‘export’
  - Implementation can be hidden
  - Order not significant

# Modules Support

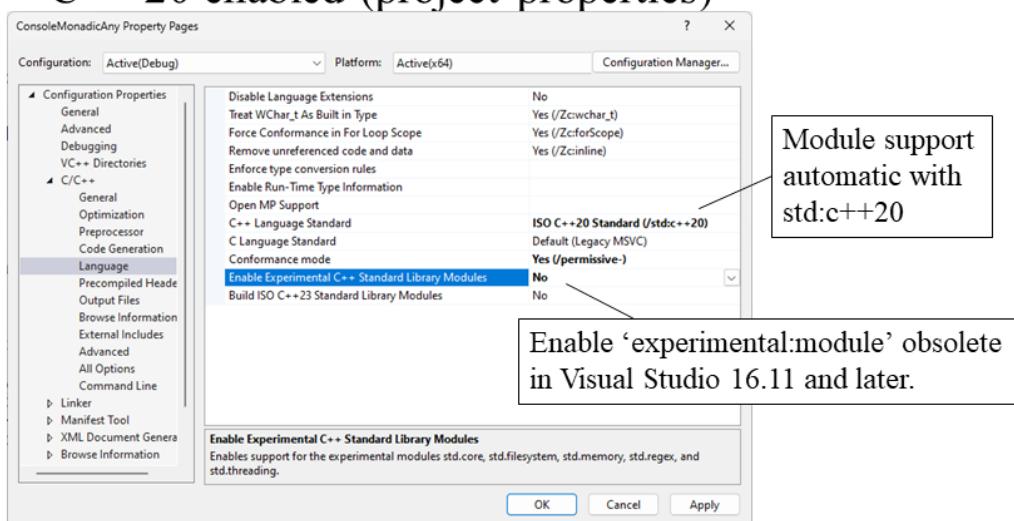
- Typically post standardisation it takes time for compiler vendors to implement all of the language features
- Modules is a feature taking time!
- Microsoft one of the first to implement Modules feature
- Also module support implemented differently by different compiler vendors
  - Different compiler options!

# Module Interface

- Define a module interface and only export what is expected to be used externally
  - ‘module’ and ‘import’ are identifiers
- Import appropriate ‘exported’ modules
  - Only processed once
  - Improved build process/reduced build time
- Import standard library header files
  - The settings required for this are compiler specific

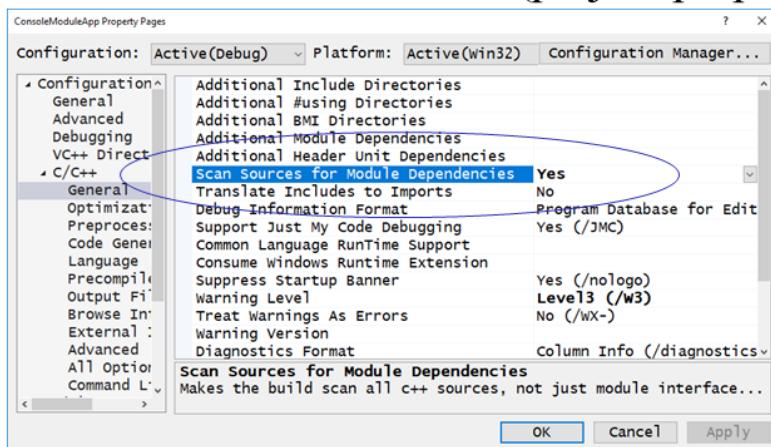
# Visual Studio Support

- Settings for modules support:
  - C++ 20 enabled (project properties)



# Visual Studio Support

- Settings required for supporting modules:
  - Scan Source for Module (project properties)



# Module Interface Unit

- A common file extension for the Module Interface Unit is ‘.ixx’
  - Single Module Interface Unit per module
  - Can have multiple Module Implementation Units
- Visual Studio projects simply add a Module
  - Provided settings correct, project should build!
- For clang compiler
  - build the ‘.ixx’ file in addition to the ‘.cpp’ file
    - Creates a ‘.ifc’ file in addition to ‘.obj’ file

# Export module

- Modules can be defined as ‘.ixx’ files:

```
module; Global module fragment after 'module'
export module flighthandling.passengers; Export 'module'.
import <string>;
export namespace flighthandling {
 class PassengerDetails { Export namespace
 std::string _name;
 int _weight;
 public:
 PassengerDetails(const std::string& name, const int& weight) :
 _name(name), _weight(weight) {}
 std::string get_name()const { return _name; }
 int get_weight()const { return _weight; }
 ...
 }
}
```

# Previous Module Explained

- The use of ‘module’ initially defines a ‘global module fragment’
  - Allows including header files for use within the module
- The export of the module name defines how this module may be imported
  - Import other modules
  - Export namespaces/classes/functions
    - But not with static linkage

## Modules continued...

- Many modules will need to include header files, but as mention, if supported these ‘header’ files may be ‘imported’ instead
- Many modules will need to import other modules, however these modules may also be exported, e.g.:

```
export import flighthandling.flights;
```

## Module Implementation Unit

- All of the implementation could be within the Module Interface Unit, however Implementation Units can be added
  - Simply files with the extension ‘.cpp’
  - Use ‘module’ keyword with module name to be implemented, but no exports
  - Implement function signatures declared within Module Interface Unit

# Implementation Unit Example

- Splitting implementation into separate file:  
`Simple.ixx`

```
export module Simple;
```

```
export inline int factorial(int a);
```

`SimpleImpl.cpp`

```
import Simple;
```

```
int main()
```

```
{
```

```
 std::cout << factorial(5) << std::endl;
```

```
}
```

```
module Simple;
```

```
int factorial(int a){
```

```
 return a == 0 ? 1 : a * factorial(a - 1);
```

```
}
```

# Implementation Units and Private

- One of the advantages of creating implementation files is that a change to this code would not require recompiling the Interface Unit
- When using just single file Module Interface Units a similar effect can be achieved by creating a private module fragment:

```
module : private;
```

# Private Module Fragment

- A Single Module Interface Unit with Private Module Fragment:

```
export module Simple;

export int cube(int a);

module :private;

int cube(int a) { return a * a * a; }

import Simple;
int main()
{
 std::cout << cube(5) << std::endl;
}
```

# Modules Partitions

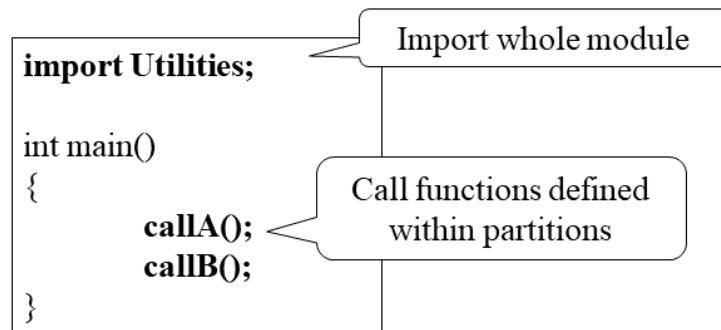
- Modules may have partitions
  - Single partition within a ‘module unit’

```
module;
#include <iostream>
export module Utilities:A; Export Module Partition ‘A’
export void do_workA() { std::cout << "From A..." << std::endl; }
```

```
export module Utilities; Export Module
import :A; Import Partition ‘A’
import :B;
export void callA() { do_workA(); }
export void callB() { do_workB(); }
```

# Using Module with Partitions

- To use the module defined in partitions simply import the module:



# Reusing Modules

- To reuse the Modules these can be package in static libraries
- Using Microsoft tools a reference can be added to reference the library
- To create the library it needs to be compile with appropriate options

# Static Library (Modules)

- Using Visual Studio:
  - Add a new Static Library project
  - To this project add a new item:
    - Module Interface Unit (.ixx)
    - Name this appropriate
    - Add class/functions for export
  - Edit the project properties
    - C/C++/Advanced/'Compile As'
      - Compile as C++ Module Code (/interface)

# Appendix 2

## Coroutines

# Coroutines

- This section gives an overview to this new language feature:
  - Coroutine Introduction
  - Generator
  - Experimental Generator
  - Coroutines and the Stack
  - Subroutines and Coroutines
  - Experimental Libraries
  - Task and `co_return`

## Coroutines Introduction

- The section is only an overview of Coroutines
- C++ 20 provides the language features however the library support will not be provided until C++ 23
- This feature allows asynchronous programming within C++
- Similar or related features have been supported in some languages/frameworks for some time, i.e. C#/.NET and Python

# Coroutines - continued

- Coroutines are functions which contain at least one of the new keyword `co_yield`, `co_await`, `co_return`
- Whilst the language features have been added, the library features have not
  - Currently only provided in experimental form, such as (with Visual Studio):
    - `std::experimental::generator`

# Generator

- An example of a coroutine is that of a function using a ‘generator’
  - The motivation for this is to allow ‘lazy’ generation of data, rather than require all of the data to be ‘created’/‘returned’ in one go as below:

```
std::vector<int> sequence(int start, int end) {
 std::vector<int> result(end - start);
 std::iota(result.begin(), result.end(), 1);
 return result;
}
```

Vector populated with increasing sequence

Whole vector returned

# Usage of ‘sequence’ function

- The sequence function could be used in a number of ways, depending on the API to be used, such as:

```
std::vector<int> seq = sequence(1, 20);

for (auto& n : seq) {
 std::cout << n << ", ";
}
std::cout << std::endl;
```

# Motivation

- The previous example of the ‘sequence’ function returned a vector
- Two particular problems arise:
  - It is of a predetermined fixed size
  - The whole vector is returned
- What if the size was not known or the sequence did not have a predetermined end or was ‘infinite’?

# Experimental Generator

- Using the experimental Generator (Visual Studio) (available in C++23), the sequence function could be written as:

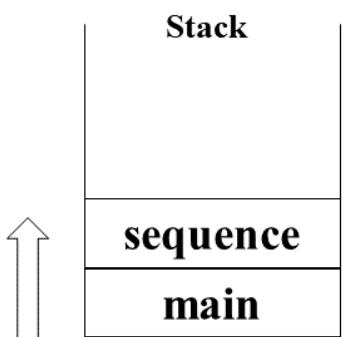
```
#include <experimental/generator>
std::experimental::generator<int> myco_sequence(int start, int end) {
 while (start != end)
 co_yield start++;
}
```

- And used:

```
for (auto& n : myco_sequence(1,20)){
 std::cout << n << ", ";
}
std::cout << std::endl;
```

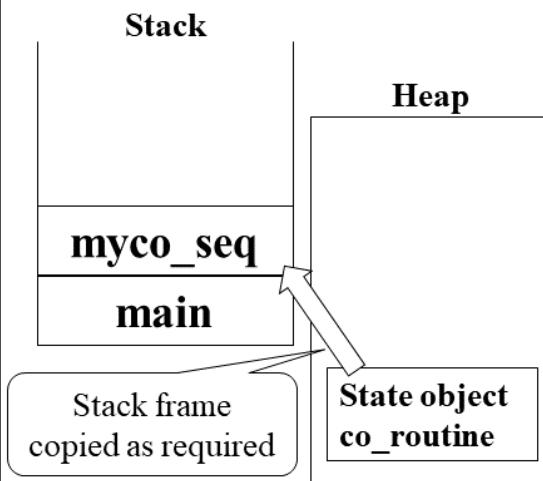
# Coroutines and the Stack

## Conventional Function Call



Stack grows as functions called  
and shrinks as functions return

## Coroutine Function



# Subroutines and Coroutines

- So far the term function has been used, however the more general term subroutine is sometimes used
  - Subroutines allow:
    - Call
    - Return
  - Coroutines additionally allow:
    - Suspend
    - Resume
- Coroutine is a generalisation of a subroutine!

## Coroutines keywords and Restrictions

- The meaning of the various keywords:

| Keyword   | Description                    |
|-----------|--------------------------------|
| co_await  | Suspends and returns control   |
| co_return | Return statement for coroutine |
| co_yield  | Returns value and suspends     |

- Restrictions:

- Cannot use variadic arguments, conventional ‘return’ or return auto (or concept)
- Coroutines cannot be consteval, constexpr, constructors, destructors or main!

# Unhandled Exceptions

- Care needs to be taken to avoid unhandled exceptions within coroutines
  - An unhandled exception within the a coroutine is caught resulting in a call to `promise.unhandled_exception()`
  - `promise.final_suspend()` is called
  - Attempting to resume this coroutine results in undefined behaviour

# Experimental Libraries

- The current language support is intended for library writers, but there are already a number of libraries which can be experimented with
  - Not recommended for release applications at this time (but useful to gain some familiarity)
  - One illustration is from Reza Arjmandi
    - [https://github.com/reza-arjmandi/cpp20\\_coroutines](https://github.com/reza-arjmandi/cpp20_coroutines)
  - This later library also illustrates the use of the some other features

## task and co\_return

- Provided a suitable library implementation exist (C++23) the following form of code should be possible:

```
task<int> make_answer(){
 co_return 42;
}

task<std::string> make_question_answer(){
 auto task = make_answer();
 auto answer = co_await task;
 co_return "The answer: " + std::to_string(answer);
}
```

## promise\_type Implementation

- Implementing the ‘generate’ or ‘task’ class is not straightforward
  - Requires a nested type ‘promise\_type’ with a number of member functions:

| Member function     |
|---------------------|
| get_return_object   |
| initial_suspend     |
| final_suspend       |
| yield_value         |
| unhandled_exception |